



Image Credit: [Tiia Monto](#)^[1]

DSTA Coursework 1: Dataset dimensionality analysis

Montel Moore | MSc Data Science | 08/02/2021

PHASE 1: INTRODUCTION




Figure 1: Bike sharing in London over time.[2]

Active travel is becoming increasingly relevant as we shift to a more health and planet conscious world. In 2010, a cycle hire scheme was implemented in London, and it saw significant uptake over the decade.

This report analyses a Kaggle dataset built around the hourly bike hires in London, its data management side and the dimensionalities that it involves.

Phase 2

2.1 DESCRIPTION AND SOURCE OF DIMENSIONS

This dataset is a time series, focused on the first two dimensions: *timestamp* and count (*cnt*). Each row is an instance of data corresponding to the timestamp, and the dataset consists of hourly timestamps from the dates of 01/04/15 – 01/03/17. The count is the number of bike rentals or “shares”, of London’s Santander cycles per hour, collected from Transport for London [3].

The dataset also contains weather data dimensions (temperature (*t1*), temperature feels like (*t2*), humidity percentage (*hum*), *wind_speed* and *weather_code*) from [freemeteo.com](#) [4]. Bank holidays were identified from the UK government website [5], and this Boolean dimension *is_holiday* was included.

Finally, *season* and the Boolean dimension *is_weekend* were added. *Season* and *is_weekend* were likely to have been inferred based on the compiler’s knowledge of the UK’s seasonal

patterns and the Georgian calendar. Figure 2 shows a preview of the dataset as a pandas DataFrame.

	timestamp	cnt	t1	t2	hum	wind_speed	weather_code	is_holiday	is_weekend	season
0	2015-01-04 00:00:00	182	3.0	2.0	93.0	6.0	3.0	0.0	1.0	3.0
1	2015-01-04 01:00:00	138	3.0	2.5	93.0	5.0	1.0	0.0	1.0	3.0
2	2015-01-04 02:00:00	134	2.5	2.5	96.5	0.0	1.0	0.0	1.0	3.0
3	2015-01-04 03:00:00	72	2.0	2.0	100.0	0.0	1.0	0.0	1.0	3.0
4	2015-01-04 04:00:00	47	2.0	0.0	93.0	6.5	1.0	0.0	1.0	3.0
...
17409	2017-01-03 19:00:00	1042	5.0	1.0	81.0	19.0	3.0	0.0	0.0	3.0
17410	2017-01-03 20:00:00	541	5.0	1.0	81.0	21.0	4.0	0.0	0.0	3.0
17411	2017-01-03 21:00:00	337	5.5	1.5	78.5	24.0	4.0	0.0	0.0	3.0
17412	2017-01-03 22:00:00	224	5.5	1.5	76.0	23.0	4.0	0.0	0.0	3.0
17413	2017-01-03 23:00:00	139	5.0	1.0	76.0	22.0	2.0	0.0	0.0	3.0

Figure 2: First five and last five rows from original Kaggle dataset

The Kaggle challenge is “to predict the future bike shares”. As of 08/02/2021, 38 notebooks have been submitted for predicting the future demand of bike shares.

Phase 3

Dimensional analysis: write down the main aggregate measures of the dataset: number of data points, number of dimensions. Select a small number of dimensions that you consider the key to understanding how data is distributed. Describe and comment those dimensions (e.g., range of the dimension, quality of the data, possible data quality/integrity issues) in your essay.

3.1 DATA CLEANING/PRE-PROCESSING

The original dataset contains two years of sequential hourly data (from 04/01/2015 to 01/03/2017), organised in 17414 rows and 10 columns/dimensions as described in Phase 2. Pre-processing techniques outlined below were performed in order to make the provide more information and to make the dataset suitable for statistical modelling and analysis.

3.1.1 Addition of variables

The *Timestamp* variable was further discretised into variables representing the *day-of-week* and *hour* (out of 23). Another dimension (*observation*) was also produced. *Observation* is the number of hours from the starting time (e.g. 04/01/2015 00:00:00 = 0, 04/01/2015 01:00:00 = 1); this makes it convenient to model the effect of time on bike shares.

The *weather_code*, *season*, *day-of-week* and *hour* variables were one-hot dummy encoded to enable a correlation analysis.

An additional dimension, *precipitation(mm)* was added to the dataset. This was deemed necessary because rainfall can affect ridership [6]. The “weather code” categories 7 and 10 describe rainfall but provide ambiguity regarding the level of rainfall. The precipitation data was gathered from CEDA.

3.1.2 Missing data

The dataset was expected to have 17544 rows worth of data (731 days * 24 hours per day), however it only contained 17414 rows of data. A list of the full 17544 hours was generated separately from the dataset, and the dataset was compared against this list in order to select a sample of missing hours. The “missing hours” were searched for in the original weather and cycle count data sources; it was found that the weather data had some missing observations, which were the likely cause for the “missing hours”.

Figure 3 shows the number of data points per hourly interval. Some hours have more data points than others because the dataset author chose to omit the timestamps with missing weather data. Due to the large volume of data and the relatively even distribution of hours, it is unlikely that this will have a significant biasing effect on predictive models built from the data.




Figure 3: Total number of data points per hourly interval

3.1.3 Asynchronous time intervals

The historical weather data from freemeteo.com is recorded at the 20th and 50th minutes per hour, as shown in the “time” column of Figure 4. This does not align with the timestamps in the dataset, which were recorded every 0th minute of the hour. This means that the data set author interpolated the weather data match to every 0th minute of the hour. This was likely done by either taking the 20th or 50th minute data alone, or by averaging both values for each hour. Weather is unlikely to change dramatically within the space of 30 minutes, therefore this discrepancy is unlikely to have a significant biasing effect on predictive models built from the data.






Time	Temperature	Relative Temperature	Wind	Wind Gust	Rel. humidity	Dew Point	Pressure	Icon	Description
00:20	-1°C	-6°C	 17 Km/h	N/A	75%	-5°C	1003.0mb		Rain
00:50	-1°C	-6°C	Variable at 17 Km/h	N/A	75%	-5°C	1003.0mb		Rain
01:20	-1°C	-6°C	 17 Km/h	N/A	75%	-5°C	1003.0mb		Rain

Figure 4: Sample of historical weather data captured from freemeteo.com

3.2 ANALYSIS OF CRITICAL FEATURES

3.2.1 Feature importance and correlation scoring

Two methods were used to evaluate the importance each feature with respect to the dependent variable, bike shares: Pearson’s correlation coefficient and MSE-based feature

importance's extracted from a random-trees regressor model. The Pearson's correlation coefficient is the linear correlation between the dependent variable and the independent variable. The random-forest regressor was used in addition to Pearson's because the Pearson's correlation coefficient is weak at modelling non-linear relationships. In addition, random-forests models remove some multi-collinearity, while Pearson's does not take it into account. There was a high degree of collinearity between independent predictors t1 and t2 ($r = 0.988344$), therefore t2 was removed from the feature analysis. Though the random-forest model takes collinearity into account, it still ranked t1 and t2 similarly because each individual tree would place a roughly equal weighting to t1 or t2 when deciding to split on a variable.

Table 1 and table 2 show the top three predictive variables based on correlation coefficients and random forest feature importance's. These were calculated for two datasets: one with continuous variables only, and another with mixed and continuous variables. The continuous dataset variables will be used for PCA, and the mixed variables dataset will be used for a (FAMD) factor analysis of mixed data.

Top 3 Correlation coefficients			
Continuous variables only		Mixed variables	
hum	-0.463	hum	-0.463
t1	0.389	t1	0.389
hour	0.324	08:00	0.334

Table 2: Top 3 correlation coefficients for the dataset with continuous variables only, and both mixed and continuous variables.

Top 3 Random forest feature imporances			
Continuous variables only model		Mixed variables model	
hour	0.623	hour	0.552
t1	0.125	is_weekend	0.150
observation #	0.100	t1	0.086

Table 1: Top 3 features ranked on mse-based feature importance, for the dataset with continuous variables only, and mixed and continuous variables.

3.2.2 Descriptive statistics

The following section will detail the range and quality of the most important dimensions and the appended dimension, *precipitation(mm)*.

Hour

The hour of the day is a strong predictor of the bike shares, being the most important variable for the random forest models. It appears to have a x^7 polynomial relationship to the total bike shares, and a general trend of more bike users in the later hours of the day, though there is a morning peak at 8:00 am, which was the most important hour category when hours were encoded as dummy variable. As mentioned in section 3.1.2, there is a roughly even spread of data points for each hour, which makes hour a reliable predictor.




Figure 5: Average bike shares per hour of day

Temperature

Temperature had a significant positive correlation with the number of shares, which is to be expected because cyclists are less likely to cycle in cold weather[7]. This is especially evident once the temperature exceeds 25 degrees, where it becomes unlikely for there to be less than 1000 cycle hires per hour. Temperatures ranged from -1.5 to 34, however temperatures over 28 were relatively uncommon.




Figure 8a: Scatter plot of temperature (t1) against bike shares per hour

Figure 8b: Distribution of temperature

Figure 8c: Temperature descriptive statistics

Humidity

Humidity was the most important variable according to the Pearson's correlation coefficient, however it was ranked as a less important feature for the random forest model. The correlation coefficient was negative which means that cycles were less likely to be hired as the humidity increased towards 100%. Humidity ranged from 20.5% to 100%, however humidity values below 31% were relatively uncommon.




Figure 11a: Scatter plot of humidity (hum) against bike shares per hour




Figure 11b: Distribution of humidity (hum)

	hum
count	17414
mean	72.324954
std	14.313186
min	20.5
25%	63
50%	74.5
75%	83
max	100

Figure 11c: Humidity descriptive statistics

Timestamp/observation number

The timestamp/observation number has a small, but significant and positive correlation with the number of bike shares ($p < 0.05$, $r = 0.04$). From Figure 12 it is evident that the number of bike shares follows a yearly cyclical pattern, with the number of bike shares being slightly higher in 2016 than in 2015.




Figure 12: Scatter plot of timestamp against bike shares per hour

Weekend/Weekday

Is_weekend had a significant negative correlation with the number of bike shares. This is likely due to a higher commuter volume on weekdays vs weekends. Figure 13 shows the average bike shares for the weekday vs weekend.




Figure 13: Average hourly bike shares for weekend days vs weekdays. 0.0 represents weekdays, 1.0 represents weekends

Precipitation/rainfall

The appended variable (precipitation) did not appear to have a significant positive or negative correlation with the number of bike shares per hour. Though Figure 16 shows a negative correlation between precipitation and bike shares, there were not enough recordings of precipitation > 0 to make a significant difference. In addition, there appears to be a non-linear relationship between precipitation and bike shares which cannot be captured with the Pearson's correlation coefficient.




Figure 15a: Scatter plot of precipitation against bike shares per hour




Figure 15b: Distribution of precipitation.

SUMMARY/FUTURE GOALS

The Kaggle “London” bike sharing dataset is a reliable and usable dataset which involves carefully picked dimensions. In part two of this report, Principal Component Analysis (PCA) will be investigated to reduce the dimensions into a set of eigenvectors which can be used as a predictor of the total number of bike shares. A feasibility study will be undertaken on a 2nd dimensionality reduction method, Factor Analysis of Mixed Data (FAMD).

REFERENCES

1. Monto, T., 2018. *Santander bicycles on the Exhibition Road in London..* [image] Available at: <https://commons.wikimedia.org/wiki/File:Santander_Cycles.jpg> [Accessed 19 February 2021].
2. Number of hires of Santander bikes from June 2011 to May 2019 (annual rolling average).. [image] Available at: <https://commons.wikimedia.org/wiki/File:Santander_bikes.png> [Accessed 19 February 2021].
3. *cycling.data.tfl.gov.uk.* [online] Available at: <<https://cycling.data.tfl.gov.uk/>> [Accessed 19 February 2021].
4. Freemeteo.co.uk. n.d. *The Weather.* [online] Available at: <<https://freemeteo.co.uk/>> [Accessed 19 February 2021].
5. GOV.UK. n.d. *UK bank holidays.* [online] Available at: <<https://www.gov.uk/bank-holidays>> [Accessed 19 February 2021].
6. Ahmed, Farhana & Rose, G. & Jacob, C.. (2010). *Impact of weather on commuter cyclist behaviour and implications for climate change adaptation.* ATRF 2010: 33rd Australasian Transport Research Forum.
7. Tin Tin, S., Woodward, A., Robinson, E. and Ameratunga, S., 2012. Temporal, seasonal and weather effects on cycle volume: an ecological study. *Environmental Health*, 11(1).

APPENDIX

Code for DSTA CW1

February 20, 2021

1 Phase 1

We have chosen to select the “London Bike Sharing” dataset

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.stats import ttest_ind
import statsmodels.api as sm
import matplotlib.ticker as plticker
import datetime
import locale
locale.setlocale(locale.LC_ALL, '')
```

```
[1]: 'English_United Kingdom.1252'
```

```
[2]: bike_shares_data = pd.read_csv('data\london_merged.csv', sep = ",")
```

```
[3]: bike_shares_data
```

```
[3]:
```

	timestamp	cnt	t1	t2	hum	wind_speed	weather_code	\
0	2015-01-04 00:00:00	182	3.0	2.0	93.0	6.0	3.0	
1	2015-01-04 01:00:00	138	3.0	2.5	93.0	5.0	1.0	
2	2015-01-04 02:00:00	134	2.5	2.5	96.5	0.0	1.0	
3	2015-01-04 03:00:00	72	2.0	2.0	100.0	0.0	1.0	
4	2015-01-04 04:00:00	47	2.0	0.0	93.0	6.5	1.0	
...	
17409	2017-01-03 19:00:00	1042	5.0	1.0	81.0	19.0	3.0	
17410	2017-01-03 20:00:00	541	5.0	1.0	81.0	21.0	4.0	
17411	2017-01-03 21:00:00	337	5.5	1.5	78.5	24.0	4.0	
17412	2017-01-03 22:00:00	224	5.5	1.5	76.0	23.0	4.0	
17413	2017-01-03 23:00:00	139	5.0	1.0	76.0	22.0	2.0	

	is_holiday	is_weekend	season
0	0.0	1.0	3.0
1	0.0	1.0	3.0
2	0.0	1.0	3.0

3	0.0	1.0	3.0
4	0.0	1.0	3.0
...
17409	0.0	0.0	3.0
17410	0.0	0.0	3.0
17411	0.0	0.0	3.0
17412	0.0	0.0	3.0
17413	0.0	0.0	3.0

[17414 rows x 10 columns]

```
[4]: bike_shares_data.reset_index(level=0, inplace=True)
```

```
[5]: bike_shares_data = bike_shares_data.rename({"index": "observation"}, axis = 1)
```

```
[6]: bike_shares_data
```

```
[6]:
```

	observation	timestamp	cnt	t1	t2	hum	wind_speed	\
0	0	2015-01-04 00:00:00	182	3.0	2.0	93.0	6.0	
1	1	2015-01-04 01:00:00	138	3.0	2.5	93.0	5.0	
2	2	2015-01-04 02:00:00	134	2.5	2.5	96.5	0.0	
3	3	2015-01-04 03:00:00	72	2.0	2.0	100.0	0.0	
4	4	2015-01-04 04:00:00	47	2.0	0.0	93.0	6.5	
...	
17409	17409	2017-01-03 19:00:00	1042	5.0	1.0	81.0	19.0	
17410	17410	2017-01-03 20:00:00	541	5.0	1.0	81.0	21.0	
17411	17411	2017-01-03 21:00:00	337	5.5	1.5	78.5	24.0	
17412	17412	2017-01-03 22:00:00	224	5.5	1.5	76.0	23.0	
17413	17413	2017-01-03 23:00:00	139	5.0	1.0	76.0	22.0	

	weather_code	is_holiday	is_weekend	season
0	3.0	0.0	1.0	3.0
1	1.0	0.0	1.0	3.0
2	1.0	0.0	1.0	3.0
3	1.0	0.0	1.0	3.0
4	1.0	0.0	1.0	3.0
...
17409	3.0	0.0	0.0	3.0
17410	4.0	0.0	0.0	3.0
17411	4.0	0.0	0.0	3.0
17412	4.0	0.0	0.0	3.0
17413	2.0	0.0	0.0	3.0

[17414 rows x 11 columns]

```
[7]: bike_shares_data['timestamp'] = pd.to_datetime(bike_shares_data['timestamp'])
```

```
[8]: bike_shares_data["season"].replace(to_replace = 0.0, value = "spring", inplace_
    ↳= True)
bike_shares_data["season"].replace(to_replace = 1.0, value = "summer", inplace_
    ↳= True)
bike_shares_data["season"].replace(to_replace = 2.0, value = "fall", inplace =_
    ↳True)
bike_shares_data["season"].replace(to_replace = 3.0, value = "winter", inplace_
    ↳= True)
```

```
[9]: l = bike_shares_data.shape[0]
bike_shares_data.insert(2, "hour", [datetime.datetime.
    ↳strftime(bike_shares_data["timestamp"][i], '%H') for i in range(l)])
bike_shares_data.insert(2, "day of week", [datetime.datetime.
    ↳strftime(bike_shares_data["timestamp"][i], '%A') for i in range(l)])
bike_shares_data["hour"] = bike_shares_data["hour"].astype("int64")
```

```
[10]: rainfall_data = pd.read_csv("data\Rainfall_2015.csv", sep = ",").append(
    pd.read_csv("data\Rainfall_2016.csv", sep = ",").append(pd.
    ↳read_csv("data\Rainfall_2017.csv", sep = ","), ignore_index =_
    ↳True), ignore_index = True)
rainfall_data
```

```
[10]:
```

	ob_end_time	prcp_amt
0	04/01/2015 00:00	0.0
1	04/01/2015 01:00	0.0
2	04/01/2015 02:00	0.0
3	04/01/2015 03:00	0.0
4	04/01/2015 04:00	0.0
...
17537	03/01/2017 19:00	0.0
17538	03/01/2017 20:00	0.0
17539	03/01/2017 21:00	0.0
17540	03/01/2017 22:00	0.0
17541	03/01/2017 23:00	0.0

[17542 rows x 2 columns]

```
[11]: for i in range(rainfall_data.shape[0]):
    rainfall_data.iloc[i,0] = datetime.datetime.strptime(rainfall_data.
    ↳iloc[i,0], "%d/%m/%Y %H:%M")
```

```
[12]: rainfall_data = rainfall_data.rename(columns={"ob_end_time": "timestamp",_
    ↳"prcp_amt": "precipitation(mm)"})
```

```
[13]: bike_shares_data["precipitation(mm)"] = 0.0
```

```
[14]: for i in range(bike_shares_data["timestamp"].shape[0]):
        a = rainfall_data[rainfall_data["timestamp"] ==
        bike_shares_data["timestamp"].iat[i]]["precipitation(mm)"]
        if a.size == 1:
            bike_shares_data["precipitation(mm)"].iat[i] = a
```

2 Plot average, or total share count per hour - average likely better

```
[15]: hours = [str(i) + ":00" for i in bike_shares_data["hour"].unique()]
col = ["cnt"] + hours
hour_of_day_dummy = pd.DataFrame(data = bike_shares_data["cnt"], columns = col)
hour_of_day_dummy = hour_of_day_dummy.fillna(0)
for i in range(bike_shares_data.shape[0]):
    hour = bike_shares_data["hour"][i]
    hour_of_day_dummy[str(hour) + ":00"][i] = 1
hour_of_day_dummy
```

```
[15]:
```


	cnt	0:00	1:00	2:00	3:00	4:00	5:00	6:00	7:00	8:00	...	14:00	\
0	182	1	0	0	0	0	0	0	0	0	...	0	
1	138	0	1	0	0	0	0	0	0	0	...	0	
2	134	0	0	1	0	0	0	0	0	0	...	0	
3	72	0	0	0	1	0	0	0	0	0	...	0	
4	47	0	0	0	0	1	0	0	0	0	...	0	
...	
17409	1042	0	0	0	0	0	0	0	0	0	...	0	
17410	541	0	0	0	0	0	0	0	0	0	...	0	
17411	337	0	0	0	0	0	0	0	0	0	...	0	
17412	224	0	0	0	0	0	0	0	0	0	...	0	
17413	139	0	0	0	0	0	0	0	0	0	...	0	
...	
17409		15:00	16:00	17:00	18:00	19:00	20:00	21:00	22:00	23:00			
0	0	0	0	0	0	0	0	0	0	0			
1	0	0	0	0	0	0	0	0	0	0			
2	0	0	0	0	0	0	0	0	0	0			
3	0	0	0	0	0	0	0	0	0	0			
4	0	0	0	0	0	0	0	0	0	0			
...			
17409	0	0	0	0	1	0	0	0	0	0			
17410	0	0	0	0	0	1	0	0	0	0			
17411	0	0	0	0	0	0	1	0	0	0			
17412	0	0	0	0	0	0	0	1	0	0			
17413	0	0	0	0	0	0	0	0	0	1			

[17414 rows x 25 columns]


```
[16]: df1 = bike_shares_data.groupby("hour").mean()
ax = sns.barplot(x = df1.index, y = "cnt", data = df1)
ax.set(xlabel = "Hour of day", ylabel = "Average bike shares", title = "Average bike shares per hour of day")

plt.show()

#constant added for intercept,
linear_model=sm.OLS(hour_of_day_dummy["cnt"], sm.add_constant(hour_of_day_dummy.
    drop(labels = ["cnt", "0:00"], axis = 1)))
result=linear_model.fit()
print(result.summary())
```



OLS Regression Results

Dep. Variable:	cnt	R-squared:	0.605
Model:	OLS	Adj. R-squared:	0.604
Method:	Least Squares	F-statistic:	1156.
Date:	Sat, 20 Feb 2021	Prob (F-statistic):	0.00
Time:	00:24:05	Log-Likelihood:	-1.3834e+05
No. Observations:	17414	AIC:	2.767e+05
Df Residuals:	17390	BIC:	2.769e+05
Df Model:	23		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	290.6091	25.373	11.453	0.000	240.875	340.343
1:00	-89.9779	35.883	-2.508	0.012	-160.313	-19.643
2:00	-154.3054	35.921	-4.296	0.000	-224.713	-83.898
3:00	-196.3636	35.921	-5.467	0.000	-266.771	-125.956
4:00	-217.2957	35.921	-6.049	0.000	-287.704	-146.888
5:00	-179.9018	35.921	-5.008	0.000	-250.310	-109.494
6:00	176.0176	35.859	4.909	0.000	105.731	246.304
7:00	1178.1361	35.859	32.855	0.000	1107.850	1248.422
8:00	2592.2141	35.883	72.240	0.000	2521.879	2662.549
9:00	1362.4101	35.846	38.007	0.000	1292.148	1432.672
10:00	774.0516	35.871	21.579	0.000	703.741	844.362
11:00	860.6096	35.846	24.008	0.000	790.347	930.872
12:00	1143.1083	35.822	31.911	0.000	1072.894	1213.322
13:00	1215.1299	35.834	33.910	0.000	1144.892	1285.368
14:00	1181.3950	35.834	32.969	0.000	1111.157	1251.633
15:00	1274.0123	35.822	35.565	0.000	1203.798	1344.226
16:00	1579.1923	35.809	44.100	0.000	1509.002	1649.382
17:00	2538.9760	35.834	70.854	0.000	2468.738	2609.214
18:00	2338.4348	35.834	65.258	0.000	2268.197	2408.673
19:00	1360.9012	35.846	37.965	0.000	1290.639	1431.163
20:00	769.1722	35.846	21.458	0.000	698.910	839.434
21:00	450.6402	35.859	12.567	0.000	380.354	520.926
22:00	301.8005	35.871	8.414	0.000	231.490	372.111
23:00	149.0432	35.908	4.151	0.000	78.660	219.427
=====						
Omnibus:		1579.612	Durbin-Watson:			0.420
Prob(Omnibus):		0.000	Jarque-Bera (JB):			10454.999
Skew:		-0.141	Prob(JB):			0.00
Kurtosis:		6.785	Cond. No.			25.0
=====						


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[17]: df1 = bike_shares_data.groupby("hour").mean()
ax = sns.histplot(x = bike_shares_data["hour"], bins=np.arange(25)-0.5)

ax.set(xlabel = "Hour", ylabel = "Count", title = "Number of Data points per_
↪hourly interval")
ax.set(xticks=range(0, 24), xticklabels=list(range(0,24)))

plt.show()
```



```
[18]: pd.DataFrame(bike_shares_data["hour"].value_counts().sort_index())
```

```
[18]:
```

	hour
0	724
1	724
2	721
3	721
4	721
5	721
6	726
7	726
8	724
9	727
10	725
11	727
12	729
13	728
14	728
15	729
16	730
17	728
18	728
19	727

```
20 727
21 726
22 725
23 722
```


3 Plot average share count per day-of-week in histogram

```
[19]: days = [i for i in bike_shares_data["day of week"].unique()]
      col = ["cnt"] + days
      day_of_week_dummy = pd.DataFrame(data = bike_shares_data["cnt"], columns = col)
      day_of_week_dummy = day_of_week_dummy.fillna(0)
      for i in range(bike_shares_data.shape[0]):
          day = bike_shares_data["day of week"][i]
          day_of_week_dummy[day][i] = 1

[20]: df1 = bike_shares_data.groupby("day of week").mean()
      ax = sns.barplot(x = df1.index, y = "cnt", data = df1)
      ax.set(xlabel = "Day of Week", ylabel = "Average bike shares", title = "Average bike shares per day of week")

      plt.show()

      #constant added for intercept,
      linear_model=sm.OLS(day_of_week_dummy["cnt"], sm.add_constant(day_of_week_dummy.
          ↪drop(labels = ["cnt", "Friday"], axis = 1)))
      result=linear_model.fit()
      print(result.summary())
```



OLS Regression Results

```

=====
Dep. Variable:          cnt      R-squared:                0.011
Model:                  OLS      Adj. R-squared:           0.010
Method:                 Least Squares      F-statistic:           31.49
Date:                   Sat, 20 Feb 2021    Prob (F-statistic):      6.88e-38
Time:                   00:24:07           Log-Likelihood:         -1.4633e+05
No. Observations:       17414             AIC:                   2.927e+05
Df Residuals:           17407             BIC:                   2.927e+05
Df Model:                6
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	1182.7727	21.808	54.235	0.000	1140.026	1225.519
Sunday	-223.2054	30.672	-7.277	0.000	-283.325	-163.086
Monday	-52.5019	30.663	-1.712	0.087	-112.604	7.600
Tuesday	47.3327	30.672	1.543	0.123	-12.787	107.452
Wednesday	61.6363	30.720	2.006	0.045	1.421	121.851
Thursday	76.0379	30.711	2.476	0.013	15.841	136.235
Saturday	-187.2189	30.795	-6.080	0.000	-247.579	-126.859

```

=====
Omnibus:                 3519.550      Durbin-Watson:           0.442
Prob(Omnibus):           0.000        Jarque-Bera (JB):        6330.073
=====

```

Skew:	1.292	Prob(JB):	0.00
Kurtosis:	4.431	Cond. No.	7.93

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[21]: np.asarray(bike_shares_data["cnt"])
```

```
[21]: array([182, 138, 134, ..., 337, 224, 139], dtype=int64)
```


4 Plot t1 “actual temperature” against count

```
[22]: ax = sns.scatterplot(x = bike_shares_data["t1"], y = bike_shares_data["cnt"])

ax.set(xlabel = "Temperature", ylabel = "Count")

plt.show()

linear_model=sm.OLS(bike_shares_data["cnt"], sm.
    ↳add_constant(bike_shares_data["t1"]))
result=linear_model.fit()
print(result.summary())
```



OLS Regression Results

```

=====
Dep. Variable:          cnt    R-squared:                0.151
Model:                  OLS    Adj. R-squared:            0.151
Method:                 Least Squares    F-statistic:        3101.
Date:                  Sat, 20 Feb 2021    Prob (F-statistic):    0.00
Time:                  00:24:07    Log-Likelihood:       -1.4500e+05
No. Observations:      17414    AIC:                 2.900e+05
Df Residuals:          17412    BIC:                 2.900e+05
Df Model:               1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	199.0395	18.569	10.719	0.000	162.641	235.437
t1	75.7183	1.360	55.685	0.000	73.053	78.384

```

=====
Omnibus:                 3716.877    Durbin-Watson:           0.508
Prob(Omnibus):            0.000    Jarque-Bera (JB):        7309.695
Skew:                     1.294    Prob(JB):                 0.00
Kurtosis:                 4.837    Cond. No.                 33.6
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


5 Plot t2 “feels like temperature” against count

```

[23]: sns.scatterplot(x = bike_shares_data["t2"], y = bike_shares_data["cnt"])
      plt.show()

      linear_model=sm.OLS(bike_shares_data["cnt"], sm.
      ↪add_constant(bike_shares_data["t2"]))
      result=linear_model.fit()
      print(result.summary())

```

OLS Regression Results

```

=====
Dep. Variable:          cnt      R-squared:                0.136
Model:                  OLS      Adj. R-squared:           0.136
Method:                 Least Squares      F-statistic:         2745.
Date:                   Sat, 20 Feb 2021    Prob (F-statistic):      0.00
Time:                   00:24:07    Log-Likelihood:         -1.4515e+05
No. Observations:       17414      AIC:                   2.903e+05
Df Residuals:           17412      BIC:                   2.903e+05
Df Model:                1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	445.6968	15.349	29.038	0.000	415.612	475.782
t2	60.5342	1.155	52.394	0.000	58.270	62.799

```

=====
Omnibus:                 3625.855    Durbin-Watson:           0.501
Prob(Omnibus):           0.000      Jarque-Bera (JB):        6962.646
Skew:                    1.278      Prob(JB):                 0.00
Kurtosis:                4.750      Cond. No.                 26.8
=====

```

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly

specified.

6 Plot humidity against count

```
[24]: sns.scatterplot(x = bike_shares_data["hum"], y = bike_shares_data["cnt"])
plt.show()

linear_model=sm.OLS(bike_shares_data["cnt"], sm.
    ↪add_constant(bike_shares_data["hum"]))
result=linear_model.fit()
print(result.summary())
```



OLS Regression Results

```
=====
Dep. Variable:          cnt      R-squared:                0.214
Model:                  OLS      Adj. R-squared:           0.214
Method:                 Least Squares      F-statistic:          4748.
Date:                  Sat, 20 Feb 2021     Prob (F-statistic):       0.00
Time:                  00:24:07      Log-Likelihood:         -1.4432e+05
No. Observations:      17414      AIC:                   2.887e+05
Df Residuals:          17412      BIC:                   2.887e+05
Df Model:               1
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
-----	-----	-----	-----	-----	-----	-----
const	3681.2262	37.547	98.043	0.000	3607.630	3754.822
hum	-35.0933	0.509	-68.909	0.000	-36.092	-34.095
=====	=====	=====	=====	=====	=====	=====
Omnibus:	4870.089		Durbin-Watson:		0.555	
Prob(Omnibus):	0.000		Jarque-Bera (JB):		12270.925	
Skew:	1.539		Prob(JB):		0.00	
Kurtosis:	5.727		Cond. No.		380.	
=====	=====	=====	=====	=====	=====	=====


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

7 Plot wind speed against count

```
[25]: sns.scatterplot(x = bike_shares_data["wind_speed"], y = bike_shares_data["cnt"])
plt.show()

linear_model=sm.OLS(bike_shares_data["cnt"], sm.
    ↳add_constant(bike_shares_data["wind_speed"]))
result=linear_model.fit()
print(result.summary())
```



OLS Regression Results

```

=====
Dep. Variable:          cnt      R-squared:                0.014
Model:                  OLS      Adj. R-squared:           0.013
Method:                 Least Squares      F-statistic:         238.7
Date:                  Sat, 20 Feb 2021     Prob (F-statistic):    1.70e-53
Time:                  00:24:08      Log-Likelihood:       -1.4630e+05
No. Observations:      17414      AIC:                  2.926e+05
Df Residuals:          17412      BIC:                  2.926e+05
Df Model:               1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	888.7350	18.378	48.359	0.000	852.713	924.757
wind_speed	15.9848	1.035	15.451	0.000	13.957	18.013

```

=====
Omnibus:                 3811.201      Durbin-Watson:           0.443
Prob(Omnibus):            0.000      Jarque-Bera (JB):        7293.776
Skew:                     1.350      Prob(JB):                 0.00
Kurtosis:                 4.662      Cond. No.                 40.1
=====

```

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

8 Plot precipitation(mm) against count

```

[26]: sns.scatterplot(x = bike_shares_data["precipitation(mm)"], y =
      ↪bike_shares_data["cnt"])
plt.show()
linear_model=sm.OLS(bike_shares_data["cnt"], sm.
      ↪add_constant(bike_shares_data["precipitation(mm)"]))
result=linear_model.fit()
print(result.summary())

```



OLS Regression Results

```

=====
Dep. Variable:          cnt      R-squared:                0.009
Model:                  OLS      Adj. R-squared:            0.009
Method:                 Least Squares      F-statistic:          153.4
Date:                  Sat, 20 Feb 2021      Prob (F-statistic):    4.32e-35
Time:                  00:24:08      Log-Likelihood:        -1.4635e+05
No. Observations:      17414      AIC:                  2.927e+05
Df Residuals:          17412      BIC:                  2.927e+05
Df Model:               1
Covariance Type:       nonrobust
=====

```

```

=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
const          1160.6244      8.308     139.692     0.000     1144.339
1176.910
precipitation(mm) -263.7346     21.291    -12.387     0.000    -305.467
-222.002
=====

```

```

=====
Omnibus:          3653.241      Durbin-Watson:           0.445
Prob(Omnibus):    0.000      Jarque-Bera (JB):        6764.134
Skew:             1.318      Prob(JB):                0.00
=====

```

Kurtosis: 4.542 Cond. No. 2.61

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[27]: bike_shares_data

```
[27]:      observation      timestamp day of week  hour  cnt  t1  t2  \
0          0 2015-01-04 00:00:00    Sunday    0  182  3.0  2.0
1          1 2015-01-04 01:00:00    Sunday    1  138  3.0  2.5
2          2 2015-01-04 02:00:00    Sunday    2  134  2.5  2.5
3          3 2015-01-04 03:00:00    Sunday    3   72  2.0  2.0
4          4 2015-01-04 04:00:00    Sunday    4   47  2.0  0.0
...
17409      17409 2017-01-03 19:00:00    Tuesday   19 1042  5.0  1.0
17410      17410 2017-01-03 20:00:00    Tuesday   20  541  5.0  1.0
17411      17411 2017-01-03 21:00:00    Tuesday   21  337  5.5  1.5
17412      17412 2017-01-03 22:00:00    Tuesday   22  224  5.5  1.5
17413      17413 2017-01-03 23:00:00    Tuesday   23  139  5.0  1.0

      hum  wind_speed  weather_code  is_holiday  is_weekend  season  \
0      93.0         6.0           3.0         0.0          1.0  winter
1      93.0         5.0           1.0         0.0          1.0  winter
2      96.5         0.0           1.0         0.0          1.0  winter
3     100.0         0.0           1.0         0.0          1.0  winter
4      93.0         6.5           1.0         0.0          1.0  winter
...
17409     81.0        19.0           3.0         0.0          0.0  winter
17410     81.0        21.0           4.0         0.0          0.0  winter
17411     78.5        24.0           4.0         0.0          0.0  winter
17412     76.0        23.0           4.0         0.0          0.0  winter
17413     76.0        22.0           2.0         0.0          0.0  winter

      precipitation(mm)
0              0.0
1              0.0
2              0.0
3              0.0
4              0.0
...
17409          0.0
17410          0.0
17411          0.0
17412          0.0
17413          0.0
```

[17414 rows x 14 columns]

9 Histogram of weather code against count

Weather code likely not the best sort of predictor

```
[28]: weather = ["weather code " + str(int(i)) for i in np.
↳ sort(bike_shares_data["weather_code"].unique())

col = ["cnt"] + weather

col
```

```
[28]: ['cnt',
      'weather code 1',
      'weather code 2',
      'weather code 3',
      'weather code 4',
      'weather code 7',
      'weather code 10',
      'weather code 26']
```

```
[29]: weather_code_dummy = pd.DataFrame(data = bike_shares_data["cnt"], columns = col)
weather_code_dummy = weather_code_dummy.fillna(0)
weather_code_dummy

for i in range(bike_shares_data.shape[0]):
    code = "weather code " + str(int(bike_shares_data["weather_code"][i]))
    weather_code_dummy[code][i] = 1

weather_code_dummy
```

```
[29]:
```

	cnt	weather code 1	weather code 2	weather code 3	weather code 4 \
0	182	0	0	1	0
1	138	1	0	0	0
2	134	1	0	0	0
3	72	1	0	0	0
4	47	1	0	0	0
...
17409	1042	0	0	1	0
17410	541	0	0	0	1
17411	337	0	0	0	1
17412	224	0	0	0	1
17413	139	0	1	0	0
		weather code 7	weather code 10	weather code 26	

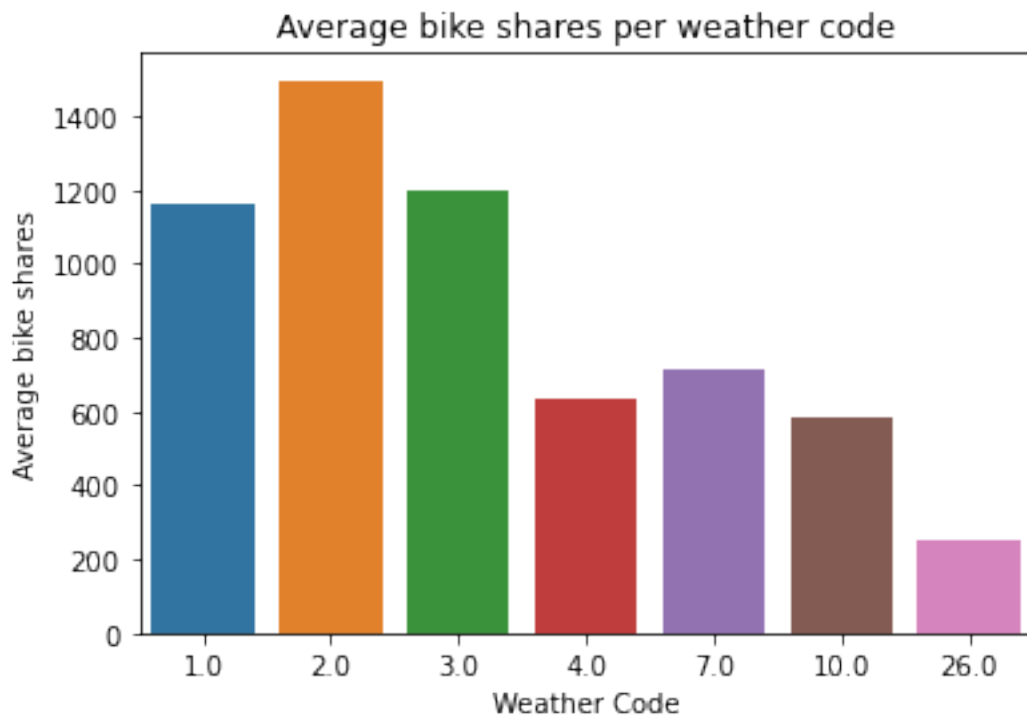
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0
...
17409	0	0	0
17410	0	0	0
17411	0	0	0
17412	0	0	0
17413	0	0	0

[17414 rows x 8 columns]

```
[30]: df1 = bike_shares_data.groupby("weather_code").mean()
ax = sns.barplot(x = df1.index, y = "cnt", data = df1)
ax.set(xlabel = "Weather Code", ylabel = "Average bike shares", title = "Average bike shares per weather code")

plt.show()

linear_model=sm.OLS(weather_code_dummy["cnt"], sm.
    ↳add_constant(weather_code_dummy.drop(labels = ["cnt", "weather code 2"],
    ↳axis = 1)))
result=linear_model.fit()
print(result.summary())
```



OLS Regression Results

```

=====
Dep. Variable:          cnt      R-squared:                0.065
Model:                  OLS      Adj. R-squared:           0.065
Method:                 Least Squares      F-statistic:         203.0
Date:                   Sat, 20 Feb 2021    Prob (F-statistic):    4.09e-251
Time:                   00:24:09    Log-Likelihood:       -1.4583e+05
No. Observations:      17414      AIC:                  2.917e+05
Df Residuals:          17407      BIC:                  2.917e+05
Df Model:               6
Covariance Type:       nonrobust
=====

```

```

===
               coef      std err          t      P>|t|      [0.025
0.975]
-----
---
const          1496.1775      16.519      90.571      0.000      1463.798
1528.557
weather code 1  -334.0885      21.258     -15.716      0.000     -375.756
-292.421
weather code 3  -301.0530      24.143     -12.469      0.000     -348.377
-253.730
weather code 4  -860.9466      32.013     -26.894      0.000     -923.696
-798.198
weather code 7  -783.2111      28.055     -27.917      0.000     -838.201
-728.221
weather code 10 -912.7489      280.900      -3.249      0.001    -1463.342
-362.156
weather code 26 -1245.3275      136.457      -9.126      0.000    -1512.796
-977.859
=====

```

```

=====
Omnibus:                 3749.423      Durbin-Watson:           0.510
Prob(Omnibus):            0.000      Jarque-Bera (JB):        7230.494
Skew:                     1.320      Prob(JB):                 0.00
Kurtosis:                 4.730      Cond. No.                 38.8
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

10 Average count for “is holiday = TRUE” vs average count for “is holiday = False”

May need to control for other variables as holidays are few and far between


```
[31]: df1 = bike_shares_data.groupby("is_holiday").mean()
ax = sns.barplot(x = df1.index, y = "cnt", data = df1)
ax.set(xlabel = "Is_holiday", ylabel = "Average bike shares", title = "Average_
    ↳bike shares for public holidays")

plt.show()

linear_model=sm.OLS(bike_shares_data["cnt"], sm.
    ↳add_constant(bike_shares_data["is_holiday"]))
result=linear_model.fit()
print(result.summary())

df2 = bike_shares_data.filter(items = ["is_holiday", "cnt"], axis = 1)
ax = sns.boxplot(x = "is_holiday", y = "cnt", data = df2)
ax.set(xlabel = "Is_holiday", ylabel = "Average bike shares", title = "Average_
    ↳bike shares for public holidays")

plt.show()
```



OLS Regression Results

```

=====
Dep. Variable:          cnt      R-squared:                0.003
Model:                  OLS      Adj. R-squared:           0.003
Method:                 Least Squares      F-statistic:         46.66
Date:                  Sat, 20 Feb 2021     Prob (F-statistic):    8.71e-12
Time:                  00:24:10      Log-Likelihood:       -1.4640e+05
No. Observations:      17414      AIC:                 2.928e+05
Df Residuals:          17412      BIC:                 2.928e+05
Df Model:               1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	1151.5252	8.304	138.668	0.000	1135.248	1167.802
is_holiday	-381.9991	55.922	-6.831	0.000	-491.611	-272.387

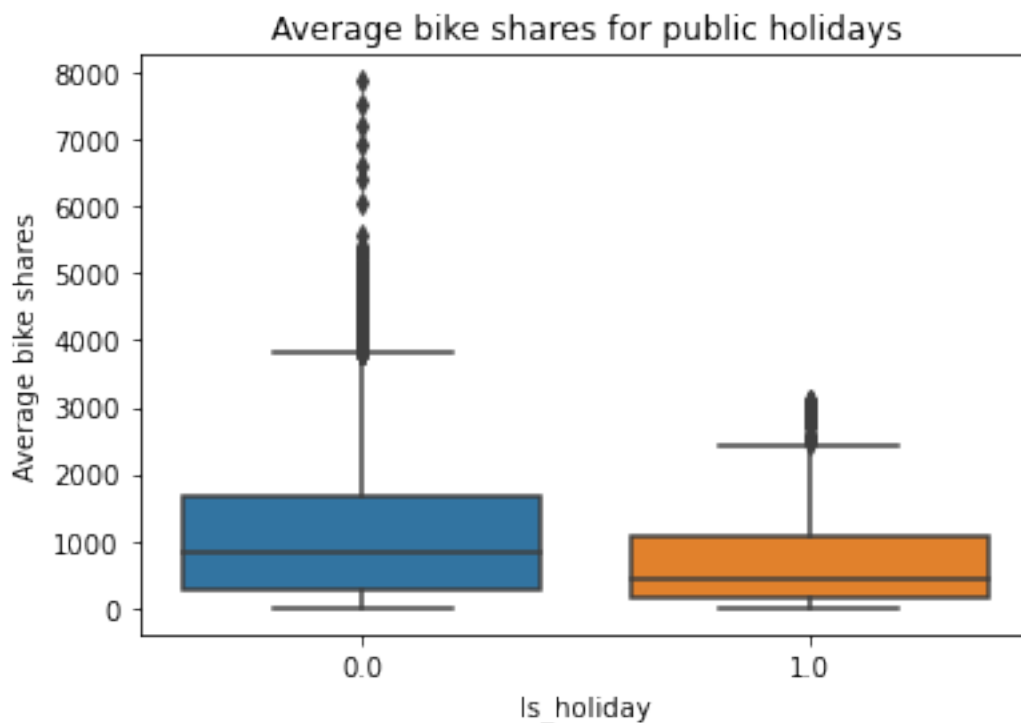
```

=====
Omnibus:                 3661.265      Durbin-Watson:           0.438
Prob(Omnibus):           0.000      Jarque-Bera (JB):       6776.708
Skew:                    1.321      Prob(JB):               0.00
Kurtosis:                4.535      Cond. No.               6.81
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.




11 Average count for “is weekend = TRUE” vs average count for “is weekend = False”

```
[32]: df1 = bike_shares_data.groupby("is_weekend").mean()
ax = sns.barplot(x = df1.index, y = "cnt", data = df1)
ax.set(xlabel = "Is_weekend", ylabel = "Average bike shares", title = "Average_
↳bike shares for weekday vs weekend")

plt.show()

linear_model=sm.OLS(bike_shares_data["cnt"], sm.
↳add_constant(bike_shares_data["is_weekend"]))
result=linear_model.fit()
print(result.summary())
```



OLS Regression Results

```
=====
Dep. Variable:          cnt    R-squared:                0.009
Model:                  OLS    Adj. R-squared:           0.009
Method:                 Least Squares    F-statistic:         163.7
```

```

Date:                Sat, 20 Feb 2021    Prob (F-statistic):        2.63e-37
Time:                00:24:10           Log-Likelihood:          -1.4634e+05
No. Observations:    17414              AIC:                   2.927e+05
Df Residuals:        17412              BIC:                   2.927e+05
Df Model:            1
Covariance Type:     nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const         1209.2748      9.682     124.897      0.000     1190.297     1228.253
is_weekend   -231.8591     18.124     -12.793      0.000     -267.383     -196.335
=====

Omnibus:                 3542.728    Durbin-Watson:           0.442
Prob(Omnibus):            0.000    Jarque-Bera (JB):        6404.084
Skew:                    1.296    Prob(JB):                0.00
Kurtosis:                4.451    Cond. No.                2.44
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

12 Average count for season

```

[33]: season = [str(i) for i in np.sort(bike_shares_data["season"].unique())]

col = ["cnt"] + season

col

```

```

[33]: ['cnt', 'fall', 'spring', 'summer', 'winter']

```

```

[34]: bike_shares_data["season"]

```

```

[34]: 0      winter
      1      winter
      2      winter
      3      winter
      4      winter
      ...
     17409   winter
     17410   winter
     17411   winter
     17412   winter
     17413   winter
Name: season, Length: 17414, dtype: object

```

```
[35]: season_dummy = pd.DataFrame(data = bike_shares_data["cnt"], columns = col)
season_dummy = season_dummy.fillna(0)
season_dummy

for i in range(bike_shares_data.shape[0]):
    code = bike_shares_data["season"][i]
    season_dummy[code][i] = 1

season_dummy
```


```
[35]:      cnt  fall  spring  summer  winter
0      182     0      0      0      1
1      138     0      0      0      1
2      134     0      0      0      1
3       72     0      0      0      1
4       47     0      0      0      1
...
17409  1042     0      0      0      1
17410   541     0      0      0      1
17411   337     0      0      0      1
17412   224     0      0      0      1
17413   139     0      0      0      1
```

[17414 rows x 5 columns]

```
[36]: df1 = bike_shares_data.groupby("season").mean()
ax = sns.barplot(x = df1.index, y = "cnt", data = df1)
ax.set(xlabel = "Season", ylabel = "Average bike shares", title = "Average bike_
    ↪ shares per season")

plt.show()

linear_model=sm.OLS(season_dummy["cnt"], sm.add_constant(season_dummy.
    ↪ drop(labels = ["cnt", "summer"], axis = 1)))
result=linear_model.fit()
print(result.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          cnt      R-squared:                0.045
Model:                  OLS      Adj. R-squared:           0.044
Method:                 Least Squares      F-statistic:           270.3
Date:                   Sat, 20 Feb 2021    Prob (F-statistic):      1.66e-171
Time:                   00:24:11    Log-Likelihood:          -1.4603e+05
No. Observations:       17414      AIC:                    2.921e+05
Df Residuals:           17410      BIC:                    2.921e+05
Df Model:                3
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	1464.4652	16.015	91.441	0.000	1433.073	1495.857
fall	-285.5110	22.760	-12.545	0.000	-330.122	-240.900
spring	-360.6336	22.640	-15.929	0.000	-405.011	-316.256
winter	-642.7361	22.724	-28.285	0.000	-687.277	-598.195

```


=====
Omnibus:                 3321.878    Durbin-Watson:           0.457
Prob(Omnibus):           0.000      Jarque-Bera (JB):        5850.871
Skew:                    1.232      Prob(JB):                 0.00
Kurtosis:                 4.412      Cond. No.                 4.78
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[37]: import matplotlib.ticker as ticker
ax = sns.scatterplot(x = bike_shares_data["timestamp"], y = bike_shares_data["cnt"])
ax.set(xlabel = "Timestamp", ylabel = "Count")
ax.xaxis.set_major_locator(ticker.LinearLocator(6))
plt.show()
linear_model=sm.OLS(bike_shares_data["cnt"], sm.add_constant(bike_shares_data["observation"]))
result=linear_model.fit()
print(result.summary())
```



OLS Regression Results

```
=====
Dep. Variable:          cnt      R-squared:                0.002
Model:                  OLS      Adj. R-squared:           0.002
Method:                 Least Squares      F-statistic:          28.02
Date:                   Sat, 20 Feb 2021    Prob (F-statistic):      1.21e-07
Time:                   00:24:11            Log-Likelihood:         -1.4641e+05
No. Observations:       17414              AIC:                  2.928e+05
Df Residuals:           17412              BIC:                  2.928e+05
=====
```

```

Df Model:                                1
Covariance Type:                        nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const          1067.7680      16.432      64.980      0.000      1035.559      1099.977
observation      0.0087       0.002       5.294      0.000         0.005         0.012
=====
Omnibus:                 3658.931   Durbin-Watson:                 0.438
Prob(Omnibus):             0.000   Jarque-Bera (JB):             6768.183
Skew:                      1.321   Prob(JB):                      0.00
Kurtosis:                  4.533   Cond. No.                      2.01e+04
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.01e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
[38]: corr_matrix_continuous = bike_shares_data.drop(labels = ["timestamp", "season",
    ↳ "day of week", "is_weekend", "is_holiday", "weather_code"], axis = 1).corr()
```

```
[39]: corr_matrix_appended = pd.concat((bike_shares_data, day_of_week_dummy.
    ↳ drop(labels = "cnt", axis = 1),
        weather_code_dummy.drop(labels = "cnt", axis = 1),
        season_dummy.drop(labels = "cnt", axis = 1),
        hour_of_day_dummy.drop(labels = "cnt", axis = 1)), axis = 1, join =
    ↳ "outer").corr()
count_corr = corr_matrix_appended.filter(items = ["cnt"], axis = 1).abs().
    ↳ sort_values(by = "cnt", axis = 0, ascending = False)
```

```
[40]: corr_matrix_continuous
```

```
[40]:
```

	observation	hour	cnt	t1	t2 \
observation	1.000000	0.001678	0.040086	0.132239	0.143846
hour	0.001678	1.000000	0.324423	0.168708	0.153956
cnt	0.040086	0.324423	1.000000	0.388798	0.369035
t1	0.132239	0.168708	0.388798	1.000000	0.988344
t2	0.143846	0.153956	0.369035	0.988344	1.000000
hum	0.119287	-0.295653	-0.462901	-0.447781	-0.403495
wind_speed	-0.126083	0.141792	0.116295	0.145471	0.088409
precipitation(mm)	0.025859	-0.004743	-0.093463	-0.004250	-0.007846

	hum	wind_speed	precipitation(mm)
observation	0.119287	-0.126083	0.025859
hour	-0.295653	0.141792	-0.004743

cnt	-0.462901	0.116295	-0.093463
t1	-0.447781	0.145471	-0.004250
t2	-0.403495	0.088409	-0.007846
hum	1.000000	-0.287789	0.153130
wind_speed	-0.287789	1.000000	0.043441
precipitation(mm)	0.153130	0.043441	1.000000

```
[41]: corr_matrix_continuous.filter(items = ["cnt"], axis = 1).abs().sort_values(by =  
      ↪ "cnt", axis = 0, ascending = False)
```

```
[41]:
```

	cnt
cnt	1.000000
hum	0.462901
t1	0.388798
t2	0.369035
hour	0.324423
wind_speed	0.116295
precipitation(mm)	0.093463
observation	0.040086

```
[42]: count_corr
```

```
[42]:
```

	cnt
cnt	1.000000
hum	0.462901
t1	0.388798
t2	0.369035
8:00	0.333934
17:00	0.324647
hour	0.324423
18:00	0.286043
4:00	0.204898
3:00	0.200889
5:00	0.197736
2:00	0.192833
1:00	0.180904
weather code 2	0.178668
summer	0.171869
winter	0.170381
weather_code	0.166633
0:00	0.163633
weather code 7	0.148419
weather code 4	0.141802
16:00	0.140090
23:00	0.134830
6:00	0.130034
wind_speed	0.116295

22:00	0.105780
9:00	0.098088
19:00	0.097798
is_weekend	0.096499
precipitation(mm)	0.093463
15:00	0.081200
21:00	0.077245
13:00	0.069807
Sunday	0.069332
14:00	0.063313
7:00	0.062596
12:00	0.055983
Saturday	0.055217
is_holiday	0.051698
weather code 26	0.048351
Thursday	0.043578
observation	0.040086
Wednesday	0.038127
Tuesday	0.032867
weather code 3	0.024265
spring	0.021024
fall	0.018929
20:00	0.016028
10:00	0.015067
Friday	0.014794
weather code 10	0.014631
weather code 1	0.012930
Monday	0.004850
11:00	0.001561

```
[43]: from sklearn.ensemble import RandomForestRegressor
a = pd.concat((bike_shares_data, day_of_week_dummy.drop(labels = "cnt", axis =
↳1),
           weather_code_dummy.drop(labels = "cnt", axis = 1),
           season_dummy.drop(labels = "cnt", axis = 1),
           hour_of_day_dummy.drop(labels = "cnt", axis = 1)), axis = 1, join =
↳"outer").drop(labels = ["timestamp", "cnt", "season", "day of week", "t2"],
↳axis = 1).values
b = pd.concat((bike_shares_data, day_of_week_dummy.drop(labels = "cnt", axis =
↳1),
           weather_code_dummy.drop(labels = "cnt", axis = 1),
           season_dummy.drop(labels = "cnt", axis = 1),
           hour_of_day_dummy.drop(labels = "cnt", axis = 1)), axis = 1, join =
↳"outer")["cnt"].values
c = RandomForestRegressor(n_estimators = 100, max_features = "auto").fit(a,b)
```

```
[44]: e = pd.concat((bike_shares_data, day_of_week_dummy.drop(labels = "cnt", axis = 1),
    ↪1),
    weather_code_dummy.drop(labels = "cnt", axis = 1),
    season_dummy.drop(labels = "cnt", axis = 1),
    hour_of_day_dummy.drop(labels = "cnt", axis = 1)), axis = 1, join = 1
    ↪"outer").drop(labels = ["timestamp", "cnt", "season", "day of week", "t2"],
    ↪axis = 1).columns
d = zip(e, c.feature_importances_)
```

```
[45]: f = list(d)
```

```
[46]: pd.DataFrame(f, columns = ["feature", "importance"]).sort_values(by = 1
    ↪"importance", axis = 0, ascending = False)
```

```
[46]:
```

	feature	importance
1	hour	0.552914
7	is_weekend	0.148761
2	t1	0.086558
35	8:00	0.058625
3	hum	0.029340
6	is_holiday	0.018605
0	observation	0.018096
46	19:00	0.011654
5	weather_code	0.011287
26	winter	0.010028
8	precipitation(mm)	0.009198
4	wind_speed	0.008393
45	18:00	0.003050
44	17:00	0.002884
14	Friday	0.002861
43	16:00	0.002743
47	20:00	0.002708
37	10:00	0.002180
20	weather code 7	0.002126
34	7:00	0.001802
38	11:00	0.001539
13	Thursday	0.001397
36	9:00	0.001362
9	Sunday	0.001251
15	Saturday	0.001244
33	6:00	0.001243
10	Monday	0.000834
23	fall	0.000812
18	weather code 3	0.000735
16	weather code 1	0.000671
24	spring	0.000667
11	Tuesday	0.000636

```

17      weather code 2      0.000620
25          summer      0.000616
12      Wednesday      0.000514
48          21:00      0.000455
50          23:00      0.000321
19      weather code 4      0.000207
39          12:00      0.000190
27          0:00      0.000175
42          15:00      0.000164
41          14:00      0.000154
40          13:00      0.000135
49          22:00      0.000121
28          1:00      0.000048
32          5:00      0.000027
29          2:00      0.000027
30          3:00      0.000013
22      weather code 26      0.000004
21      weather code 10      0.000004
31          4:00      0.000002

```

```

[47]: from sklearn.ensemble import RandomForestRegressor
a2 = bike_shares_data.drop(labels = ["timestamp", "cnt", "season", "day of week",
↳ "is_weekend", "is_holiday", "weather_code", "t2"], axis = 1).values
b2 = bike_shares_data["cnt"].values
c2 = RandomForestRegressor(n_estimators = 5, verbose = 2).fit(a2,b2)

```

```

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done   1 out of   1 | elapsed:   0.0s remaining:   0.0s

```

```

building tree 1 of 5
building tree 2 of 5
building tree 3 of 5
building tree 4 of 5
building tree 5 of 5

```

```

[Parallel(n_jobs=1)]: Done   5 out of   5 | elapsed:   0.2s finished

```

```

[48]: e2 = bike_shares_data.drop(labels = ["timestamp", "cnt", "season", "day of week",
↳ "is_weekend", "is_holiday", "weather_code", "t2"], axis = 1).columns
d2 = zip(e2, c2.feature_importances_)

```

```

[49]: f2 = list(d2)

```

```

[50]: pd.DataFrame(f2, columns = ["feature", "importance"]).sort_values(by =
↳ "importance", axis = 0, ascending = False)

```

```

[50]:          feature  importance
1          hour      0.614407
2           t1      0.133248

```

```

0      observation    0.096788
3              hum    0.076231
4      wind_speed    0.062931
5  precipitation(mm)  0.016395

```

```
[51]: bike_shares_data.filter(items = ["hour", "t1", "hum", "is_weekend",
↳ "precipitation(mm)"], axis = 1).describe()
```

```
[51]:
```

	hour	t1	hum	is_weekend \
count	17414.000000	17414.000000	17414.000000	17414.000000
mean	11.513265	12.468091	72.324954	0.285403
std	6.915893	5.571818	14.313186	0.451619
min	0.000000	-1.500000	20.500000	0.000000
25%	6.000000	8.000000	63.000000	0.000000
50%	12.000000	12.500000	74.500000	0.000000
75%	18.000000	16.000000	83.000000	1.000000
max	23.000000	34.000000	100.000000	1.000000

	precipitation(mm)
count	17414.000000
mean	0.066441
std	0.384543
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	11.800000

```
[52]: ax = sns.boxplot(bike_shares_data["t1"])


ax.set(xlabel = "Temperature")

plt.show()
```

```

C:\Users\Montel\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(


```

```
[53]: sns.boxplot(bike_shares_data["precipitation(mm)"])

plt.show()
```


C:\Users\Montel\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(



```
[54]: sns.boxplot(bike_shares_data["hum"])

plt.show()
```

```
C:\Users\Montel\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```



```
[55]: sns.scatterplot(x = bike_shares_data["hum"], y =  
    ↪bike_shares_data["precipitation(mm)"])  
  
plt.show()
```

